

Multi-Agent AI Collaboration Framework

Version 1.1 | **Date** 2026-06-12 | **Nature:** a replicable institutional design, distilled from a four-agent system in real production **New in v1.1:** Review Court & precedent system, three-tier token economy, claim/lease mechanism, long-horizon stage gates in practice, lessons expanded to eight

1. What problem does this solve

However capable a single AI assistant is, it has four structural defects:

1. **Nobody reviews it** — an AI accepting its own work lets errors accumulate silently.
2. **Context is finite** — long projects span dozens of conversations, each starting from amnesia.
3. **No governed memory** — either nothing is remembered, or everything is (including the wrong things).
4. **The human becomes the mailman** — relaying messages between AI tools by copy-paste.

This framework addresses all four with five pillars: **role separation** + **file-based coordination** + **review gates** + **governed memory** + **scheduled automation**. The whole system runs on an ordinary Windows PC — no server, no database. At its core it is **one folder and a few disciplines**.

2. Three axioms


Everything else is a corollary of these:

Axiom 1: Players don't referee. The AI that designs and reviews writes no production code; the AI that implements never accepts its own delivery. Nobody can objectively review code they wrote — role separation is the precondition for any quality loop.

Axiom 2: Files are the single source of truth. All coordination flows through plain-text files (append-only), never through anyone's conversational memory. "Where are we?" is always answered from files, never from recollection.

Axiom 3: A gate exists only where code enforces it. If "requires approval before execution" lives only in a document while a scheduler runs the armed code — the approval has already been bypassed. Every declared gate must point to the mechanism (flag, file check, permission) that enforces it.

3. Role architecture



Role	Duties	Explicitly forbidden
User	Requirements, decisions, high-risk gate approvals	Should never need to relay messages
Design/review side	Specs, root-cause diagnosis, independent review	Never edits production code
Implementation side	Builds to spec, delivers with evidence, replies item-by-item	Never self-reviews or self-passes
Mailman agent	Notifications, intake sorting, phone approval surface	Never writes to the canonical knowledge base

Roles bind to **duties, not brands** — any two or more AI tools that can read and write files qualify.

4. The coordination layer: files as the message bus

```

coordination/
├─ queue.md           ← task ledger (append-only, single source of truth)
├─ broadcast.md       ← broadcasts (append-only)
├─ README.md          ← entry format contract + claim contract
├─ staging/
│   └─ <TASK-ID>/     ← one folder per task
│       └─ spec or proposal
│       └─ reconcile (implementer's item-by-item reply)
│       └─ delivery report and evidence
│       └─ review reports
│       └─ coordination_claim.json (claim stamp)
│       └─ ...
└─ outbox/            ← durable return notices (for automated executors)

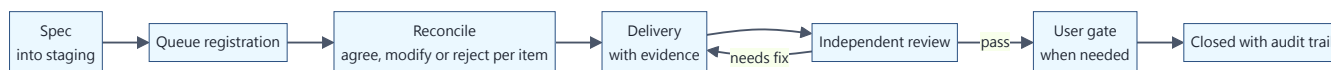
```

Core conventions:

- **Task IDs:** `YYYYMMDD-short-name`, referenced consistently throughout.
- **Append-only:** queue entries are never edited; state changes append new sections; **readers resolve each task by its LATEST section** (lesson 4).
- **Format contract:** entry fields are specified in writing, and **every automated parser must tolerate the loose formats humans actually write** — we once went blind for days because the parser only recognized the legacy dialect while every new entry used a newer one, and the health reports kept saying "all clear".
- **Claim/lease:** before processing any task folder, an agent writes a claim stamp (who, when, lease expiry); other agents back off an active claim; expired leases may be reclaimed. With

multiple agents in parallel, "two actors grabbing the same letter" becomes mechanically impossible rather than a matter of etiquette.

Task lifecycle



Reconcile is the load-bearing step: the implementer replies **agree / modify / reject for every spec item**, negotiating disagreements in place. Architectural changes require both sides' agreement before any code — preventing "the spec dreams one thing, the implementation builds another".

5. The review system

First principle of review: trust no report — re-run everything.

The reviewer's standard moves on receiving a delivery:

1. Read the diff and delivery report;
2. **Independently re-run** all verification (compile, tests, audit scripts) — never accept pasted results;
3. Run the checks the implementer forgot to list;
4. Write probes to reproduce anything suspicious;
5. Score against the fixed rubric + list required / non-blocking items;
6. Review report into staging, conclusion into the queue.

Fixed rubric (against judge drift): four dimensions at 25 points each — spec fidelity, evidence quality, safety & gate discipline, operational durability; 80 is the pass line. Every finding carries **two labels**: severity (HIGH/MEDIUM/LOW/INFO) and **certainty** (PROVEN = verified by the reviewer's own re-execution or direct inspection / INFERRED = concluded from reading / SPECULATIVE = plausible concern without direct evidence). Iron rules: **SPECULATIVE never blocks alone**; hard checks (tests, compile, audit scripts) always outrank model judgment.


Iteration is bounded: the needs_fix loop runs at most THREE rounds; a third failure stops the loop and escalates — to the user, or back to design. Never an unbounded ping-pong. Fix rounds address **ONLY** the listed required items; opportunistic refactoring is forbidden.

Two review layers: an unattended scheduled first-pass (read-only) plus an interactive deep review when a human is present. **The authority order is fixed:** hard checks > interactive tool-enabled review > automated context-limited review — a lower tier never overrides a higher one. An

automated reviewer that cannot see the evidence it needs must return "evidence gap — escalate", never needs_fix or pass (lesson 5).

6. The Review Court: reviews that learn from reviews

The review system ensures *this* review is right; the Review Court ensures the *next* one is better. Its core is **case law**:



- **A precedent** is a significant ruling distilled into a standalone entry — failure pattern, how to detect it, the ruling, applicability bounds. One-off bugs don't earn precedents; **patterns likely to recur across tasks** do.
- **Mandatory pre-read**: every reviewer (automated included) loads the accepted precedents as a checklist before reviewing.
- **Precedents are governed too**: candidate → review → acceptance, the same lifecycle as all shared memory; stale precedents are deprecated, not silently forgotten.
- Academic anchor: this institutionalizes Reflexion (memory as the cross-session outer improvement loop) — **the system's review capability rises monotonically with every review, and survives every session reset.**

Our precedent registry validated itself on opening day: hours after we recorded precedent five — "a fix written for a lesson is itself an untested path" — it caught a bug that only exploded on a refresh branch that had never once been executed.

7. The gate system

Not everything should be automated. Risk tiers:

Tier	Examples	Handling
Low	Docs, tests, internal scripts	Auto-execute, review after
Medium	Behavior changes, scheduled tasks, notification channels	Execute after review passes
High	Money, credentials, public release, data deletion	User gate: explicit human approval first

Approvals must leave structured traces. After the verbal "go ahead", record a machine-queryable field:

```
user_approved_<item>: <ISO timestamp> | <approver> | <scope>
verbatim: "OK, proceed with option A"
```


Three months later, anyone — human or AI — can answer "who approved what, when, and within what scope".

Axiom 3 in practice: we once had documentation declaring "awaiting approval" while the scheduler was already running the armed code — the system fired the first shot itself. The lesson is now institutional: **declaring a gate requires pointing at the code or config line that enforces it**, and reviews verify each such claim. The coordination contract now reads: *"a staging note, queue field, or handoff sentence is a request for a gate, not the gate itself."*

8. Governed memory (the Memory Palace)

Cross-agent long-term memory is not a log pile — it is a **shared asset requiring governance**.

```
memory/
├─ shared/
│   ├── candidates/  ← proposals (any agent may write)
│   ├── accepted/    ← canonical rules (promotion gate only)
│   ├── rejected/    ← failed proposals (kept on record)
│   ├── deprecated/  ← retired rules
│   └─ _maintenance/ ← scripts and weekly health reports
└─ private/<agent>/ ← per-agent private areas (mutually unreadable)
```




Hard rules (violations are rejected outright):

- No secrets (tokens, passwords, .env);
- No raw chat logs or session transcripts — **memory stores curated rules, not running commentary**;
- No business-domain knowledge (that belongs in the knowledge base, not agent memory);
- Every entry carries audit fields (who, when, why) + an append-only change history;
- Promotion to accepted requires the gate; no agent ever self-approves its own proposal.

Automatic aging: candidates pending past 30 days raise alerts; accepted entries carry expiry dates and return for re-review when expired and uncited. A weekly scheduled maintenance run produces health reports (pending age, expiry countdowns, duplicate IDs, secret scans).

9. The automation pipeline: making the mailman redundant



One OS scheduled task runs a pure-script pipeline every five minutes:

- **watcher** parses the queue and sorts tasks into "to implement / to reply / to approve" with priorities;
- **router** detects new deliveries in staging and generates review requests (folders without a primary artifact are marked wait-for-artifact, never churned);
- **auto reviewer** invokes a read-only AI for the first pass (claim before, release after);
- **milestone detection** pushes one-line events — "reply landed / verdict out / awaiting approval" — to the user's phone, **naming the task and whose move is next**.

Iron rules: **notification failure never fails the pipeline** (best-effort); **automation only sorts and notifies** — lifecycle changes (promotion, closure) always pass through gates; the **five-minute heartbeat itself costs zero LLM tokens** (pure local scripts; the model appears only when a real review request exists).

10. Token economy: the three-tier ladder

A personal budget is not a corporate unlimited plan — but **savings must never cost quality**. Three tiers; only the top one is the goal:

Tier	Meaning	Stance
Starvation	Feeding less, or broken, input to save money	Forbidden. Empirically: reviews fed truncated evidence issued three wrong verdicts in a row, and the repair cost more than one properly-fed review — starvation loses twice
Compression	Same judgment, fewer tokens	A legitimate saving, with a hard line: judgments made from condensed material (e.g. the precedent digest) must match judgments made from the full text — fail that equivalence check and the condensed version may not be used
Refinement	Same tokens, more judgment	The direction. The digest is re-distilled at every update — synthesizing across cases, sharpening detection cues, retiring redundancy. Judgment-per-token rises with accumulation


The floor is "quality never drops"; there is no ceiling — the system only rewards going up.

Operational principles:

- **Deterministic first:** compile/tests/audit scripts cost zero tokens and outrank the model in reliability — a red light closes the case without a drop of model ink;
- **Never re-do identical work:** unchanged artifact hashes re-emit the previous verdict;
- **Fold calls:** new extraction needs ride existing LLM calls as structured output — no new always-on calls;
- **Every new recurring LLM step must declare a weekly token estimate** — silent recurring spend is a high-severity finding (budgets, like gates, must live where the code runs);
- **Quality wins conflicts:** when economy and quality collide in a specific case, quality wins and the cost is recorded for redesign — never silently absorbed by degrading the review.

11. Long-horizon goal mode: stage gates

For multi-day objectives, use the implementer's long-horizon mode (e.g. Codex `/goal`) — but **autonomy is not exemption from review**:



- The goal definition is a file (objective, verifiable completion condition, constraints, per-stage scope); the goal command points at it;
- **Stage gates are file states**: the next stage opens only when a review file with a pass verdict exists in staging — mechanism, not honor system; when interactive and automated reviews conflict, interactive wins;
- Blocked over 24 hours: update the queue and pause — never skip a gate;
- The final completion condition is **an audit script** emitting proven / gated / missing per requirement; only all-proven completes the goal.

Field note: these gates passed their trial by fire unattended, overnight. The implementing AI received a *generic* review-pass document, judged on its own that it did not satisfy the stage-specific gate, and refused to advance; its attempt to summon a review directly was blocked by its own safety policy — **it did not bypass**, but filed a formal delegation, wrote full diagnostics, and stayed paused until morning. Gates prove their worth precisely when nobody is watching.

12. Eight lessons, all paid for

#	Lesson	Incident
1	Parse liberally, write conservatively	The queue evolved two dialects; automation recognized only the old one and went task-blind for days — while health reports said "all clear"
2	"Never happened" must be verified, not presumed	The notification feature had been live for days; archaeology showed it had never actually sent one production message — even its test was a dry-run
3	Gates in code; documents don't count	Docs said "awaiting approval" while the scheduler ran the armed code and fired the first shot itself
4	Append-only needs latest-state-wins reads	Tasks accumulated status sections; the scanner read only the first and resurrected long-closed tasks as zombies
5	An evidence gap is not a defect	Automated reviews kept failing deliveries they could not see; the correct ruling is "escalate to a tier with access", never a content verdict on invisibility
6	In parallel lanes, check the envelope first	With concurrent tasks, agents grabbing the wrong folder is a real risk; the cure is task-ID discipline plus the claim mechanism
7	When the reviewed party edits the review tooling, flag it	The implementer patched the reviewer itself to unblock a gate — the direction was tightening, so no harm, but "self-referential change" must be a mandatory delivery label, not a lucky catch
8	A fix written for a lesson is itself an untested path	The patch that cured scanner blindness shipped with a same-day-refresh branch that crashed the first time it ever ran — every new path needs a test that has walked it

13. Minimal viable adoption

You don't build it all at once. Four steps:

Step 1 (ten minutes): the file coordination layer. One `queue.md`, one `staging/` folder, one page of format contract. From then on, all handovers go through files; the human only says "go check your mail".

Step 2 (half a day): review discipline. Role separation, item-by-item reconcile, independent re-runs, the fixed rubric with its 80-point line, the three-round `needs_fix` cap. The quality loop now exists — without a single line of automation.

Step 3 (incremental): automate piece by piece. Scheduled heartbeat → watcher sorting → push notifications → automated first-pass review → claims → long-horizon goal mode. After each

piece, audit it against the eight lessons.

Step 4 (continuous): open the Review Court. Distill precedents from significant reviews, gate them into memory, pre-read them before every review. From this step on the system is not merely running — **it is getting stronger**. Forging is never one strike: every new discussion should put old conclusions back in the fire, fold and quench again — that is how you get Damascus steel.

There is exactly one yardstick for progress: is the human's remaining work nothing but decisions and approvals? Everything else — relaying, status-checking, chasing, forwarding — should be steadily disappearing.

Appendix: glossary

Term	Meaning
reconcile	The implementer's item-by-item reply to a spec (agree / modify / reject)
staging	The dedicated artifact folder per task
user gate	A checkpoint requiring explicit human approval
needs_fix	Review failure with a required-items list (three rounds max)
claim/lease	Claim before processing; expired leases reclaimable; prevents double work
milestone event	A state change worth notifying a human about
Memory Palace	Governed shared memory (candidates-to-accepted lifecycle)
precedent	A reusable pattern distilled from a review ruling (detection + ruling + bounds)
certainty label	A finding's evidence grade: PROVEN / INFERRED / SPECULATIVE
re-distillation	Digest re-refined at every promotion; judgment density rises with accumulation
proven / gated / missing	An audit script's three-state verdict per completion requirement

This specification is distilled from a real production system: two AI coding agents, one messaging-gateway agent, and one user, on a single Windows PC, governing the full path from requirement to delivery through plain text files. Every lesson in this document is a post-mortem of an actual incident.